

Grease: A Private Payment Channel Protocol for Monero

Grease Team

January 21, 2026

Contents

1. Introduction	4
1.1. Payment channels in Monero	4
1.1.1. Enter Grease	4
1.1.2. Why does another chain have to be involved?	5
2. Design principles	6
2.1. Anti-principles	6
3. The Grease Channel Lifecycle	7
3.1. Overall design description	7
3.2. High-level state machine	8
3.3. New Channel	9
3.3.1. Channel Id	11
3.4. Establishing the Channel	11
3.5. Wallet transaction protocol	13
3.5.1. Overview	13
3.5.1.1. Phase 1: Customer Preprocessing	13
3.5.1.2. Phase 2: Merchant Preprocessing and Partial Signing	14
3.5.1.3. Phase 3: Customer Verification and Signing	14
3.5.1.4. Phase 4: Final Verification	14
3.6. Channel Updates	16
3.6.1. The Verifiable Consecutive One-way Function (VCOF)	18
3.6.2. Grease VCOF	18
3.7. Co-operative Close	19
3.7.1. Channel Close messages	19
3.8. Channel Dispute	21
4. The Key Escrow Service (KES) Design	22
4.1. Introduction	22
4.2. Preliminaries	22
4.3. Global keys	23
4.4. Channel Encryption keys	23
4.4.1. Security properties	24
4.5. Communication with the KES	24
4.6. Channel Initialization	25
4.6.1. What the parties send to the KES	25
4.6.2. What the KES send back to parties	25
4.6.3. OpenChannel record	26
4.6.4. Failure scenarios	26
4.7. Channel Force Close	27
4.7.1. The original Monet/AuxChannel proposal	27
4.8. Force-closing in Grease	28
4.8.1. Summary	28
4.8.2. Force-close procedure	29
4.8.3. ForceCloseRequest message	29
4.8.4. PendingChannelClose message	30

4.8.5.	ClaimChannelRequest message	31
4.9.	Channel Dispute	32
4.9.1.	Force close with consensus state	33
4.9.1.1.	ConsensusCloseRequest message	33
4.9.2.	Dispute with recent state proof	34
4.9.3.	Do Nothing	34
4.10.	Channel cleanup	34
4.10.1.	After co-operative close	34
4.10.2.	Hygiene	35
4.11.	Utility functions	35
5.	Trust assumptions in the KES	36
5.1.	The problem with private state on blockchains	37
5.1.1.	A hybrid on-chain off-chain design	37
5.1.2.	Additional applications and distribution of risk	39
5.2.	Using TEEs	39
5.2.1.	The KES on the Secret network	39
5.2.1.1.	Advantages	40
5.2.1.2.	Disadvantages	40
5.2.1.3.	Intel SGX exploits	41
5.3.	A note on economic incentives for centralized KEServices.	41
5.4.	Summary	42
6.	Future extensions and known limitations	43
7.	Nomenclature	45
7.1.	Symbols	45
7.2.	Subscripts	45
	References	46
	Index of Figures	47
	Index of Tables	47
	Index of Listings	47

1. Introduction

Monero is, alongside cash, the world's most private¹⁻³, and, arguably the best currency in circulation, but the user *experience* remains less than ideal. This comment is not necessarily aimed at user *interfaces* – for example, there are Monero wallets that are very attractive and easy to use – but the fundamental design of Monero means that:

- many, especially new, users find they can make only one payment roughly every 20 minutes when their wallet holds a single spendable output (the change from the first payment typically requires about 10 confirmations before it can be spent again),
- due to the absence of scripting capabilities, use-cases that capture the public imagination, like DeFi, are not possible in vanilla Monero.

Therefore, the *experience* of using Monero tends to be one of waiting, and limited functionality.

For Monero to achieve mass adoption, it will need to find ways to:

- provide an order of magnitude *better UX* (again, not necessarily UI). Locking UTXOs after spending and block confirmation times add significant friction to Monero and is a turn-off for new users who are already unsure about how cryptocurrency works.
- provide *instant confirmations* when purchasing with Monero.
- enable *seamless point-of-sale transactions* so that using Monero for purchases feels no different to using a credit card or Venmo.
- enable DeFi for Monero. DeFi is the future of finance. The lack of permissionless access to bank-like services (loans, insurance, and investments) is a key barrier to truly democratic money.
- provide for Monero-backed and/or privacy-maximizing stable coins.

A payment-channel solution for Monero is one of the foundational requirements for achieving these goals in Monero. The other is smart contracting functionality, but that is out of scope for this project.

1.1. Payment channels in Monero

Monero's primary function is private, fungible money. This goal very likely excludes any kind of meaningful on-chain state management for Monero, since state implies heterogeneity. And heterogeneity immediately breaks fungibility. That's not to say that some hitherto undiscovered insight won't allow this in future, but for the short and medium-term at least, any kind of state management for Monero transactions or UTXOs would have to be stored off-chain.

It makes the most sense to store this off-chain state on another decentralized, private protocol. Zero-knowledge Rollup blockchains fit the bill nicely.

It's the goal of this project to marry Monero (for private money) with a ZK-rollup chain (for private state management) to create a proof-of-concept Monero payment channel for Monero.

1.1.1. Enter Grease

The Grease protocol is a new bi-directional payment channel design with unlimited lifetime for Monero. It is fully compatible with the current Monero implementation and is also fully compatible with the upcoming FCMP++ update.

Using the Grease protocol, two peers may trustlessly¹ cooperate to share, divide and reclaim a common locked amount of Monero XMR while minimizing the online transaction costs and with minimal use of outside trusted third parties.

The Grease protocol maintains all of Monero’s security. No identifiable information about the peers’ privately owned Monero wallets is shared between the peers. This means that there is no way that privacy can be compromised. Each channel lifecycle requires two Monero transactions, with effectively unlimited near-instant updates to the channel balance in between these two transactions. This dramatically improves the scalability of Monero.

The Grease protocol is based on the original AuxChannel⁴ paper and Monet⁵ protocol. Monet provides a sketch of a payment channel design but many details were omitted or elided. Grease is a full-fledged implementation with full consideration given to state management, multiple concurrent channels, peer-to-peer communication and encryption and practical performance.

Every update and the final closure of the channel require an online interaction between channel parties. In order to prevent the accidental or intentional violation of the protocol by a peer not interacting and thus jamming the channel closure, Grease introduces an external Key Escrow Service (KES). The KES needs to run on a stateful, logic- and time-aware platform. A decentralized zero-knowledge smart contract platform satisfies this requirement while also providing the privacy-focused ethos familiar to the Monero community.

1.1.2. Why does another chain have to be involved?

Offline payment channels necessarily *require* a state management mechanism. Typically, the scripting features for a given blockchain allow for this state to be managed directly. However, Monero’s primary design goals are privacy and fungibility. Attaching state to UTXOs would create a heterogeneity that threatens these goals. Fungibility is more important than features when it comes to privacy.

The state does not have to be managed on the same chain though. Any place where the state is:

- available,
- reliable,
- has trust assumptions appropriate for the value involved,

will suffice.

The initial implementation uses any Noir-compatible execution environment that supports the Barretenberg PLONK proving system, the Aztec blockchain being one candidate.

The KES acts as a third-party judge in disputes. At initialization, each peer encrypts a secret (their ω_0 offset) for the KES. If a dispute arises, the KES identifies the violating peer and releases that peer’s secret to the wronged peer. The wronged peer can use the secret to generate a valid Monero transaction and recover funds from the channel. Only valid channel states can be unilaterally closed; fabricated updates cannot be simulated. Section 4 has full details on the KES design and implementation.

¹No trust is needed between the channel parties. The degree of trust given to the escrow service is dependent on the specific KES used. Different implementations are possible carrying different trade-offs.

2. Design principles

Grease is a bidirectional two-party payment channel. This means that funds can flow in both directions, but in the vast majority of cases, funds will flow from one party (the “customer”) to the other (the “merchant”). This designation is somewhat arbitrary, either party in a channel may assume either role. Most channels will have participants that are easily identifiable in either the customer or merchant role and we use that terminology throughout this document.

Grease embraces this use case and optimizes the design and UX based on the following assumptions:

- The merchant is responsible for recording the channel state on the ZK chain.
- The merchant pays for gas fees on the ZK chain and will need to have some amount of ZK chain tokens to pay for these fees.
- The customer will need a small quantity of ZK chain tokens if they want to dispute a channel closure. In the vast majority of cases, this won’t be necessary, since funds almost always flow in one direction from the customer to the merchant. However, in instances where this is not the case, the customer is able to dispute the channel closure by watching the ZK chain and proving that the channel was closed with outdated state.
- In the vast majority of cases, the customer opens a channel with m XMR and the merchant starts with a zero XMR balance (since the merchant is providing assets or services in exchange for Monero).
- Usually, both parties mutually close the channel. Either party *may* force close the channel, and are able to claim their funds after a predetermined timeout. In this case, the forcing party is usually the merchant since they have the greater incentive to do so in the case of an abandoned channel.

2.1. Anti-principles

The following design goals are explicitly *excluded* from the Grease design:

- Multi-hop channels. Multi-hop channels are probably *possible* in Grease, but they are not a design goal.

Taking the Lightning Network as the case study, CJ argues⁶ that the vast majority of the utility of lightning is captured by bilateral channels, with a tiny fraction of the complexity.

3. The Grease Channel Lifecycle

3.1. Overall design description

Grease largely follows the Monet⁵ design, which is a payment channel protocol that uses a key escrow service (KES) to manage the funds in the channel.

A Grease payment channel is a 2-party bidirectional channel. The most common use case is in a multi-payment arrangement between a customer and a merchant, and so we will label the parties as such.

To set up a new channel, the customer and merchant agree on the funds to be locked in the channel. They're usually all provided by the customer, but it doesn't need to be. These funds are sent to a new 2-of-2 multisig wallet, which is created on the Monero blockchain for the sole purpose of serving the channel.

The idea is that a *commitment transaction*, so-called for reasons that will be made clear later, spends the funds out of the multisig wallet back to the customer and merchant can be trustlessly, securely and rapidly updated many thousands of times by the customer and merchant without having to go on-chain.

Every time the channel is updated, the customer and merchant provide signatures that *can't be used to spend the funds* out of the multisig wallet, but *prove* that they will be able to spend the funds if a small piece of missing data is provided². When the channel is closed, the customer gives the merchant that little piece of data and the funds are spent out of the multisig wallet to the customer and merchant, closing the channel.

If the merchant cheats and tries to close the channel with an outdated state, or decides not to broadcast the commitment transaction, the customer can dispute the closure of the channel with the key escrow service (KES).

The KES is responsible for arbitrating disputes. It is ideally a permissionless, decentralized private smart contract, but it can be a centralized 3rd-party service as well. It won't be called upon for the vast majority of channel instances, but its presence is mandatory to disincentivize cheating.

Note

In fact, you could run Grease without a KES, if there is a high-trust relationship between the customer and merchant.

When the 2-of-2 multisig wallet is created, both the customer and merchant encrypt their adapter signature offset to the KES.

If, say, the merchant tries to force-close a channel using an outdated state (which is itself enacting the dispute process), or refuses to publish anything at all (in which case the customer will enact the force-close process), the customer has a certain window in which it can prove to the KES that it has a valid, more recent channel state signature.

²These signatures are called adapter signatures.

In a successful dispute, either by waiting for the challenge period to end, or the KES accepts the challenge, the KES will hand over the merchant's first adapter offset. The customer will then be able to sign any transaction in the history of the channel by reconstructing the appropriate adapter signature offset, including any states that favour the customer. This is a form of punishment that should motivate parties to behave honestly.

Warning

Once a channel is closed, neither party should use the 2-of-2 multisig wallet again, since there exists another party that can immediately spend out of that wallet.

3.2. High-level state machine

On a high level, the payment channel lifecycle goes through 6 phases:

- **New** - The channel has just been created and is entering the establishment negotiation phase. Basic information is swapped in this phase, including the public keys of the peers, the nominated KES, and the initial balance. The channel id is derived from data describing the channel, including, the initial balance, public keys and closing addresses (See Section 3.3.1). If both parties are satisfied with the proposed channel parameters, the channel moves to the **Establishing** state.
- **Establishing** - The channel is being established. This phase includes the KES establishment and funding transaction. Once the KES is established and both parties have verified the funding transaction, the parties will share an **AckChannelEstablished** message. Once acknowledged, an **OnChannelEstablished** event is emitted, and the channel will move to the **Open** state.
- **Open** - The channel is open and ready for use. Any number of channel update events can occur in this phase and the channel can remain in this state indefinitely. The channel remains in this state until the channel is closed via the amicable **Closing** state or the **Disputing** state. At each update, both parties cryptographically *commit* to the new state of the channel and collaborate to produce a new, partially-signed commitment transaction.
- **Closing** - The channel is being closed. This phase includes sharing of adapter secrets and signing of the final commitment transaction. Once both parties have signed the final commitment transaction, any party will be able to broadcast it, but by convention it will be the merchant that does so.
- **Closed** - The channel is closed. The merchant **should** inform the KES of the closure so that it can clean up any state associated with the channel. A channel can reach the **Closed** state after a co-operative close, following a resolved dispute, or after several error conditions arising during the **New** and **Establishing** phases.
- **Disputing** - The channel is being disputed because someone initiated a force-close. If the local party initiated the force close, this phase includes invoking the force-close on the KES, and waiting for the dispute window to expire so that the counterparty's secret can be recovered in order to synthesize the closing transaction. If the other party initiated the force-close, we can invoke the KES to challenge the closing state, submit the correct signature offset, ω_n , or do nothing and allow the counterparty to use your ω_0 to recover any previous

channel state. The final state transition is always to the `Closed` state, only the reason can vary. See Section 4 for details.

3.3. New Channel


A new channel is established when a Merchant shares some initialization data with a Customer out-of-band.

The customer takes this data, combines it with their own information, and sends a channel proposal to the Merchant.

There are **three** half-rounds of communication in this phase³:

1. Out-of-band channel initialization data (CID) sharing from Merchant to Customer:
 - Contact information for the merchant
 - Channel seed metadata. This includes metadata so that both merchant and customer can uniquely identify the channel throughout the channel's lifetime. This includes:
 - a random nonce id,
 - the merchant's closing address,
 - the requested initial balances,
 - the merchant's public key,
 - the dispute window duration
 - Protocol-specific initialization data. This might include commitments for parameters that will be shared later, the KES public keys that can be accepted, etc.
2. New channel proposal from Customer to Merchant
3. Accepting the proposal from Merchant to Customer

³See `server.rs:customer_establish_new_channel`



3.3.1. Channel Id

The channel id is a 65 character hexadecimal string that uniquely identifies the channel. It is defined as the prefix “XGC”, followed by the first 31 bytes of the **Blake2b-512** hash of the channel metadata in hexadecimal format.

The hash is calculated from the following transcript:

- The merchant public key, P_B , 32 bytes in little-endian byte order,
- The customer public key, P_A , 32 bytes in little-endian byte order,
- The merchant initial balance in piconero, as a 64-bit unsigned integer in little-endian byte order,
- The customer initial balance in piconero, as a 64-bit unsigned integer in little-endian byte order,
- The nominated closing address of the merchant, as a Base58 string,
- The nominated closing address of the customer, as a Base58 string,
- A merchant nonce, a 64-bit little-endian unsigned integer, randomly chosen by the merchant, and
- A customer nonce, a 64-bit little-endian unsigned integer, randomly chosen by the customer


3.4. Establishing the Channel

Establishing a channel to accept payments requires the following preparatory steps:


1. Both parties collaboratively create a new shared multisig wallet to hold the channel funds.
2. Each party calculates a unique, shared secret for the channel, κ (see Section 4.4).
3. Each party watches the Monero blockchain for the funding transaction to confirm it has been included in a block.
4. Each party encrypts their initial adapter signature offset to the KES.
5. The merchant creates a new KES commitment on the ZK-chain smart contract and commits the encrypted shares to it.
6. The merchant provides a proof of the KES commitment to the customer who can verify that the KES was set up correctly.
7. Each party creates the initial proof for their initial secret (witness) and shares it with the counterparty.
8. The customer funds the multisig wallet with the agreed initial balance.
9. Once the funding transaction is confirmed, the channel is open and ready to use.

Note

No SNARKs are required for channel establishment. However, if the KES is deployed on a ZK-enabled chain, the KES proof of knowledge proofs should be calculated in a zero-knowledge manner.



Listing 2: Establishing a new Channel



Listing 3: Establishing a new Channel, continued

3.5. Wallet transaction protocol

3.5.1. Overview

Grease uses 2-of-2 multisig transactions generated using FROST (Flexible Round-Optimized Schnorr Threshold signatures) combined with adapter signatures. The roles of customer and merchant are interchangeable, but for simplicity we will refer to the customer as the party initiating the protocol.

3.5.1.1. Phase 1: Customer Preprocessing

The Customer initiates the protocol by generating preprocessing data based on transaction details. This preprocessing step is fundamental to FROST and involves:

- Creating commitments for the signing process
- Generating nonce values that will be used during signature creation
- Producing data (denoted as C^C) that encapsulates these commitments.

The Customer then transmits both the preprocessing data and the transaction details to the Merchant.

3.5.1.2. Phase 2: Merchant Preprocessing and Partial Signing

Upon receiving the Customer's data, the Merchant carries out:

1. **Generates its own preprocessing data** (C^M) using the same transaction details
2. **Creates a partial signature** on the transaction using:
 - The transaction details
 - Its own preprocessing commitments
 - This produces a partial signature, (R_0^M, s_0^M) .
3. **Adapts the signature** by converting the partial signature into an adapter signature format:
 - Produces an adapted signature tuple $(R_0^M, Q_0^M, \hat{s}_0^M)$
 - During channel initialization, it generates a random witness value (ω_0^M) that serves as a secret offset
 - During updates, this witness will be updated using the VCOF mechanism instead of being generated randomly.

The adapter signature is a cryptographic construct that allows the Merchant to create a valid-looking signature that is “locked” by a secret witness value. This signature appears complete but cannot be verified as a standard signature without knowledge of the witness.

The Merchant transmits its preprocessing data and the adapted signature back to the Customer, but notably **withholds** the witness value ω_0^M .

3.5.1.3. Phase 3: Customer Verification and Signing


The Customer now performs verification and, if successful, creates its own adapted signature:

- **Verifies the adapter signature** $(R_0^M, Q_0^M, \hat{s}_0^M)$ using the Merchant's public key
- If verification succeeds:
 - Creates its own partial signature (R_0^C, s_0^C) on the transaction
 - Adapts this signature to produce $(R_0^C, Q_0^C, \hat{s}_0^C)$ and witness ω_0^C
 - Transmits the adapted signature to the Merchant
- If verification fails:
 - Sends an error message to the Merchant
 - Aborts the protocol

3.5.1.4. Phase 4: Final Verification

The Merchant receives the Customer's adapted signature and performs the final verification:

- **Verifies the adapter signature** $(R_0^C, Q_0^C, \hat{s}_0^C)$ using the Customer's public key
- If verification succeeds:
 - Sends confirmation to the Customer
 - Both parties now hold valid adapted signatures
- If verification fails:
 - Sends an error message to the Customer
 - Aborts the protocol



Listing 4: Creating a new multisig wallet

3.6. Channel Updates

Once a channel is open the parties may transact and update the XMR balance between themselves. This is done entirely off-chain, and happens near-instantaneously.

Note

Under CLSAG, unbroadcast Monero transactions become stale after some time because of how decoy selection is mandated and thus channels need to be kept “fresh” by providing periodic zero-delta updates. This is not an issue with FCMP++.

Post-FCMP++, the signing mechanism for Monero transactions will be such that decoy selection can be deferred until channel closing⁷. This will simplify channel updates in two important ways:

- Updates will not need to query the Monero blockchain to select decoys at update time, which presents a significant performance improvement.
- Channels can stay open indefinitely, without risk of the closing transaction becoming stale.


Either party can initiate a channel update. Since it is usually the merchant, we designate them as the initiator of the update, but the labels can be switched without loss of generality.

The merchant decides on the delta value, which can be positive (customer pays merchant) or negative (merchant refunds customer). He then carries out phase 1 on the transaction protocol to create a new Monero transaction reflecting the updated balances.

1. Privately, compute ω_i using Figure 1 from the previous ω_{i-1} .
2. Generate a DLEQ proof for T_i and Q_i .
3. Generate a ZK-SNARK proof Π_i^V using Figure 2.
4. Send the update package to the peer.

Verification consists of:

1. Verify the DLEQ proof for T_i and Q_i .
2. Verify the ZK-SNARK proof Π_i^V using Figure 3.



Listing 5: Updating a Channel

Note

The actual implementation may streamline communication rounds by batching messages together. In particular, the VCOF messages and transaction protocol messages can be sent together to reduce the total number of round trips.

3.6.1. The Verifiable Consecutive One-way Function (VCOF)

The core idea behind AuxChannel is the Verifiable Consecutive One-way Function (VCOF). The VCOF is used to generate new adapter signature offsets for each channel update in that only the party who generated the new offset can compute it, but both parties can verify that it is valid and consecutive to the previous offset.

As such, the VCOF has the following security requirements:

- *validity*: It requires that for any message, the encrypted signature can be verified.
- *unforgeability*: It is hard to forge a valid encrypted signature.
- *recoverability*: Informally speaking, recoverability requires that it is easy to recover a decryption key by knowing the encrypted signature and its original signature.

These requirements are standard for any encrypted signature scheme. There are three additional requirements specific to the VCOF:

- *Consecutiveness*: requires that the decryption-encryption key-pair used in the i^{th} update of the VCOF is derived from the decryption key of the $(i - 1)^{\text{th}}$ update.
- *Consecutive verifiability*: Given two VCOF outputs and the corresponding proof, anyone can be convinced that the outputs are consecutive, meaning that the new secret is generated from applying the VCOF to the previous secret.
- *One-wayness*: given ω_i , no one can derive any ω_j , where $0 \leq j < i$. This is important because, without one-wayness, when a channel is closed co-operatively, a malicious counterparty could broadcast *any* previous channel state to the blockchain by computing a previous state.

3.6.2. Grease VCOF

The Grease VCOF makes use of a ZK-SNARK to produce the following update algorithm (`KeyUpdate` using the AuxChannel parlance):

Informal security arguments for the Grease VCOF:

- *One-wayness*: If `H2F` is selected to be a suitable one-way hash function, then the Grease `KeyUpdate` function, Figure 1, is also one-way.
- *Consecutiveness*: If Figure 1 is used to generate ω_{i+1} from ω_i , then the proof generated by Figure 2 will verify under Figure 3. This is because the prover knows ω_i and can compute all the necessary values to satisfy the constraints in Figure 2.

VcofUpdate(i, ω_i)

1 $\omega_{i+1} = \text{H2F}(\omega_i, i)$

2 return $(\omega_{i+1}, i + 1)$

Figure 1: Grease VCOF Update function

```

VcofProve( $i, \omega_i, \omega_{i+1}, T_i, T_{i+1}$ )


---


1  $\omega = \text{H2F}(\omega_i, i)$ 
2 assert  $\omega == \omega_{i+1}$ 
3  $P_1 = \omega_{i+1} \cdot \mathbb{G}_2$ 
4 assert  $T_{i+1} == P_1$ 
5  $P_2 = \omega_{i+1} \cdot \mathbb{G}_2$ 
6 assert  $T_{i+1} == P_2$ 
7 return  $\Pi_{i+1}^V$ 


---



```

Figure 2: The ZK-SNARK Grease `VCOFProve` function

```

VcofVerify( $i, T_i, T_{i-1}, \Pi_i^V$ )


---


1 verify ZK-SNARK proof  $\Pi_i^V, T_i, T_{i-1}$ 
2 verify DLEQ proof for  $T_i, Q_i$ 


---



```

Figure 3: Grease VCOF Verify function

- Consecutive verifiability: The ZK-SNARK proof produced by Figure 2 convinces any verifier running Figure 3 that ω_{i+1} was correctly derived from ω_i using Figure 1 without revealing either secret value.

3.7. Co-operative Close

Either party can initiate a co-operative close of the channel at any time by sending their latest adapter offset to the counterparty in a `RequestChannelClose` message. As an illustration, we will assume that the customer initiates the close.

When the merchant receives the offset, it verifies it, by trying to complete the closing transaction signature for the expected update count. It then broadcasts the transaction to the network. The channel is then marked as `Closed`. If the closure was successful, the merchant responds with an `ChannelCloseSuccess` message. If any errors occur, the merchant responds with a `RequestCloseFailed` message, and the channel remains `Open`. If the merchant did not provide the transaction id in its response, the customer may use the provided offset to reconstruct and broadcast the closing transaction herself.

If a party becomes unresponsive during the co-operative close process, one may initiate a force-close via the KES, as described in Section 4.

3.7.1. Channel Close messages

A party sends a `RequestChannelClose` message to initiate a co-operative channel closure.

```


pub struct RequestChannelClose {
    /// The globally unique channel id
    channel_id: ChannelId,
    /// The latest adapter offset for the initiating party
    offset: Scalar,
    /// The update count corresponding to the latest adapter offset

```

```
    update_count: u64,  
}
```

The counterparty responds with either a `RequestCloseFailed` or `ChannelCloseSuccess` message.

```
pub struct ChannelCloseSuccess {  
    /// The globally unique channel id  
    channel_id: ChannelId,  
    /// The latest adapter offset for the responding party  
    offset: Scalar,  
    /// The transaction id of the closing transaction. Optional.  
    txid: Option<TxId>,  
}  
  
pub struct RequestCloseFailed {  
    /// The globally unique channel id  
    channel_id: ChannelId,  
    /// An error code indicating the reason for the failure  
    reason: CloseFailureReason,  
}
```



Listing 6: The Channel Close sequence

3.8. Channel Dispute

In the event of a dispute, such as when one peer becomes unresponsive or attempts to close the channel with an outdated state, the aggrieved peer initiates a force-close procedure. This process leverages the KES to ensure fair resolution and allows the wronged peer to reclaim funds according to the latest agreed channel state.

The dispute process is described in detail in Section 4.8.

4. The Key Escrow Service (KES) Design

4.1. Introduction

In short, the KES will hold an encrypted secret, ω_0 , from each of 2 parties. Either party can request their counterpart's secret if certain conditions are met.

Therefore the KES needs a keypair for each channel (k_g, P_g) under which both parties can encrypt their secrets and send it to the KES for safe-keeping using `EncryptMessage` (Figure 4). If the unlock conditions are satisfied, the KES can decrypt ω_0 and send it to the claimant.

The KES only participates in the channel during the initialization, force-close and, dispute stages. It is not active in the update and co-operative close phases. The KES also has an optional cleanup step after a channel has closed. Table 1 summarizes the KES' role throughout the channel lifecycle.

STAGE	KES REQUIRED	KES ROLE
Initialization	Yes	Securely store ω_0 for each party.
Update	No	
Co-operative close	No	
Closed	Recommended	Discard ω_0 and other cleanup.
Force Close	Yes	Deliver counterparty ω_0 after dispute window closes.
Dispute	Yes	Deliver ω_0 to counterparty if dispute conditions are satisfied.

Table 1: The role of the KES during the channel lifecycle

4.2. Preliminaries

Let $\mathbb{G}_1$ and $\mathbb{G}_2$ be two additive groups of points on elliptic curves $\mathbb{E}_1$ and $\mathbb{E}_2$, respectively.

$\mathbb{E}_1$ is defined over a prime finite field $\text{GF}_1(p_1)$, with p_1 a large prime. Let N_1 be the size of the group generated by the Generator point, G_1 for group $\mathbb{G}_1$, such that $N_1 \cdot G_1 \equiv \mathcal{O}$.

$\mathbb{E}_2$ is a SNARK-friendly curve and is similarly defined.

$H_F(\dots)$ is a one-way function under the random oracle assumption that hashes the binary representation of its input to a scalar element in the group order. Domain separation and other hygiene considerations are assumed to be properly taken care of.

Similarly, $H_P(\dots)$ is a one-way function under the random oracle assumption that hashes the binary representation of its input to a group element in $\mathbb{G}$.

The KES can utilize $\mathbb{E}_1$, $\mathbb{E}_2$, or both, depending on where it's deployed and which dispute resolution option (Section 4.8) is selected. The VCOF utilizes a ZK-SNARK, and will therefore utilize $\mathbb{E}_2$.

`EncryptMessage`. $P = k \cdot G$ is the recipient's public key. m is a scalar to be encrypted

- 1 $(R, \chi) = \text{EncryptMessage}(m, P, \mathbb{G}, \mathbb{F})$
 - 2 Select random scalar element, r where $0 < r < N$.
 - 3 Calculate:
 - 4 $R = r \cdot G$
 - 5 $P_s = r \cdot P (\equiv rk \cdot G)$
 - 6 $s = H_F(P_s)$
 - 7 $\chi = (m + s) \bmod N$
 - 8 Return (R, χ) .
-

Figure 4: The `EncryptMessage` algorithm.

`DecryptMessage`. $P = k \cdot G$ is the recipient's public key. m is the encrypted scalar

- 1 $m = \text{DecryptMessage}(R, \chi, \mathbb{G}, \mathbb{F})$
 - 2 Calculate:
 - 3 $P_s = k \cdot R (\equiv rk \cdot G)$
 - 4 $s = H_F(P_s)$
 - 5 $m = (\chi - s) \bmod N$
 - 6 Return m .
-

Figure 5: The `DecryptMessage` algorithm.

In Grease, $\mathbb{E}_1$ is always Curve25519⁸, the curve behind Monero. $\mathbb{E}_2$ can be any suitable curve. For the reference implementation of Grease, $\mathbb{E}_2$ is the BabyJubJub curve.

4.3. Global keys

We denote the curve that the KES operates on (whether $\mathbb{E}_1$ or $\mathbb{E}_2$) as $\mathbb{E}_K$.

On creation, the KES creates a new, master keypair (k_K, P_K) on $\mathbb{E}_K$. P_K is made available publicly.

4.4. Channel Encryption keys

When a merchant-customer pair create a new channel, they create an ephemeral channel id secret on $\mathbb{E}_K$, κ , using Figure 6.

1. The merchant (or customer, it doesn't matter) then encrypts κ to the KES using `EncryptMessage` (see Figure 4) on $\mathbb{E}_K$.
2. The KES calculates the unique channel keypair for $\mathbb{E}_K$:

$$\begin{aligned} k_g &= \kappa \cdot k_K \\ P_g &= k_g \cdot G \end{aligned} \tag{1}$$

3. The KES shares P_g with both parties.
4. **Critical:** The KES discards κ .

ECDH-MC

- 1 $k_s = \hat{k}_a \hat{P}_b (\equiv \hat{k}_b \hat{P}_a)$ on $\mathbb{E}_K$
 - 2 $\kappa \leftarrow H_F(k_s, \text{channel_id})$ on $\mathbb{E}_K$
 - 3 $P_\kappa \leftarrow \kappa \cdot G$ on $\mathbb{E}_K$
-

Figure 6: Diffie-Hellman shared secret derivation for new channels

 Important

$\hat{k}_a$ **must** be chosen at random for *every* new channel a party creates. Failure to do so may result in loss of funds.

Even though $\hat{k}_a$ must be unique for every channel opening, clients do not have to store a large database of keys (although they can). One could use a **master key**, k_a and derive unique channel keys using a simple deterministic scheme,

$$\hat{k}_a = H_F(k_a, \text{channel_id}) \quad (2)$$

Since the channel id is already a function of the counterparty public keys, the initial balances, and a nonce, it will be unique for each channel.

4.4.1. Security properties

The channel keys have the following properties:

- If k_g becomes compromised, the channel is compromised, but other current and future channels are still safe, since the attacker does not know κ for the other channels.
- If k_K becomes compromised, channels are *still* safe, as long as the κ for any given channel is unknown. This is why it is critical to discard κ after deriving the keys. While the damage would be contained, a compromised k_K is incredibly serious and would require the KES to start afresh with a new master keypair and inform all parties to close open channels immediately.

4.5. Communication with the KES

Either, or both parties can carry out the communication duties with the KES, since none of the communication protocols with the KES require trust between the counterparties. Every protocol provides sufficient cryptographic proof that the steps were completed correctly.

A working assumption is that the Merchant typically has more reliable internet access and possesses a more powerful machine (typically a server) than the Customer (typically a mobile phone).

Thus, for the sake of brevity, it is assumed that the Merchant acts as proxy for all KES communications. But it is certainly possible for the Customer, or both parties to perform the same tasks, with the same, idempotent result.

4.6. Channel Initialization

The Merchant (M) and Customer (C) partially sign a 2-of-2 multisig transaction with an adaptor signature. Either party needs to know their counterpart's adaptor signature offset, ω_0 to complete the signature and spend the funds out of the multisig wallet, thus closing the channel.

The primary role of the KES during channel initialization is to store the encrypted ω_0 values for each party. The full procedure for calculating ω_0 is described in Section 3.4.

It is important to note that the adaptor signature offsets, ω_0^C and ω_0^M and their corresponding points, Q_0^C and Q_0^M are elements of $\mathbb{E}_1$, the Ed25519 curve. For the KES to work with them, they need to be translated onto $\mathbb{E}_K$.

The offsets are chosen such that they are valid scalars on both curves, that is

$$0 < \omega_0 < \min(N_1, N_2) \quad (3)$$

The group element on $\mathbb{E}_1$ corresponding to ω_0 is Q_0 , while the corresponding point on $\mathbb{E}_K$ is denoted T_0 . A summary of the points and their equivalents is given in Table 2.

Note

It is possible that $\mathbb{E}_K$ is also Curve25519. This does not change the requirements or procedures. However, many things are much simpler. The DLEQ proofs are trivial, for instance.

4.6.1. What the parties send to the KES

Each party computes a discrete-log equivalence proof (DLEQ), Π_0 that proves that Q_0 and T_0 are generated from the same scalar, ω_0 .

The Customer encrypts ω_{0C} to the KES using `EncryptMessage` (Figure 4) and sends the following package to the Merchant:

- the encrypted offset, χ_C
- The channel id, `id` (Section 3.3.1),
- The Customer's public key, P_A ,
- The length of the dispute window, `dw` in seconds. The default is 86,400 (24hrs),
- A payload signature signing `id`, χ_C , `dw`, and T_0^C with $\hat{k}_a$
- The DLEQ proof, Π_0^C .

The merchant validates Π_0^C and forwards the rest of the package to the KES.

The Merchant performs the analogous tasks and forwards both his and the customer's data to the KES. He also sends his DLEQ proof, Π_0^M to the Customer, who validates it.

4.6.2. What the KES send back to parties

Once the KES has received both packages, it performs the validation procedure outlines in Figure 7

Open Channel validation

- 1 Validate the payload signatures.
 - 2 Decrypt $\chi_C \rightarrow \omega_0^C$ and $\chi_M \rightarrow \omega_0^M$.
 - 3 Are dispute windows, `dw`, values equal?
 - 4 | No? **Return** `Fail( InvalidConfiguration )`
 - 5 Store an `OpenChannel` record in private/encrypted storage.
 - 6 Calculate proof-of-knowledge proofs, Γ_C and Γ_M , for ω_0^C and ω_0^M respectively, using the per-channel key k_g (not the global key k_K).
 - 7 Send the PoK proofs to the Merchant and/or Customer.
 - 8 Discard ω_0^C and ω_0^M .
-

Figure 7: The KES validates each new channel request, returning proofs of knowledge if successful.

4.6.3. `OpenChannel` record

`OpenChannel` records the state the KES needs to carry out its duties for every channel:

- The channel id, `id`
- The dispute window, in seconds (default: 86,400)
- The proofs-of-knowledge, Γ_C and Γ_M
- The merchant public key, P_B and encrypted initial offset χ^M ,
- The customer public key, P_A and encrypted initial offset χ^C ,

Caution

Don't store any other metadata, initial balances or anything else that would reduce privacy and increase attack surface.

Important

Notwithstanding other checks that are required, a party **must not** allow channel opening to proceed without verifying both PoK proofs. For the most part, the Customer can proxy communication with the KES through the merchant, but it is **recommended** that the Client at least queries the KES directly to obtain these proofs.

Note that to verify the PoKs, parties only need the KES global public key P_K , since the per-channel public key is simply $P_g = \kappa \cdot P_K$.

4.6.4. Failure scenarios

The Customer trusts the Merchant to relay messages to the KES faithfully. Let's examine scenarios where this breaks down:

- The Merchant never relays any messages: The Customer never receives a PoK for ω_0 and refuses to sign the channel opening.

- The Merchant uses a different KES that he controls: The KES cannot decrypt χ_C , since it is encrypted to the initially agreed-upon KES.
- The KES colludes with the merchant and gives him ω_0^C : **Failure**: This allows the Merchant to spend all funds out of the multisig wallet. This is a known failure mode in both Grease and AuxChannel. If the KES operates as a trusted third party, then the risk of collusion is governed by that degree of trust. It is preferable therefore to have the KES deployed as a publicly auditable contract on a Zero-Knowledge blockchain. The keys can be stored according to the procedure described above without revealing their values and collusion then requires subverting the entire blockchain.
- The Merchant follows the procedure as described, except that it withholds the PoKs from the Customer and immediately triggers a force-close: After the dispute window ends, the channel closes, allowing the return of just the original balances to each party. The Merchant is unable to construct any valid Monero transaction besides reversing the original funding transaction.
- The Merchant follows the procedure as described, except that it withholds the PoKs and goes offline indefinitely: The Customer, after some timeout period, can query the KES directly, retrieve the PoKs, and trigger a force-close herself. She can get her funds back after the dispute window closes.

SCALAR	$\mathbb{E}_1$	$\mathbb{E}_2$	LABEL
ω_0	Q_0	T_0	Initial adapter signature offset

Table 2: Equivalent curve points in grease

4.7. Channel Force Close

4.7.1. The original Monet/AuxChannel proposal

Suppose Bob (P_B) wants to close the channel at state i , but Alice (P_A) is unresponsive.

1. Bob posts a “trigger transaction” to the KES smart contract on L2. This transaction includes his decryption key k_{di}^B , the state index i , and the L2 signatures $\sigma_i'^A, \sigma_i'^B$. This starts a timer.
2. Alice has a time window to react. She has two options:
 - **If state i is the latest, correct state**: She cooperates by posting a “release transaction”, revealing her decryption key k_{di}^A . Bob can then use this key to recover Alice’s signature and close the channel on L1. This option has the same outcome as the co-operative close but with the added overhead of involving the KES and paying L2 gas fees. It is significantly cheaper for Alice and Bob to co-operate.
 - **If state i is outdated (e.g., Bob is trying to use an old state)**: Alice posts an “update transaction” on L2, where $j > i$ is the actual latest state index. This transaction proves that Bob’s claim is stale. To verify the proof, the KES is required to execute the VCOF $j - i$ times to verify Alice’s claim. Depending on the nature of the L2 this may or may not be an expensive operation.
3. **Timeout**: If Alice does not respond within the time window, the KES contract automatically releases her **initial** decryption key k_{d0}^A on L2.
4. **Resolution**:

- If Alice released k_{di}^A , Bob can close the channel at state i .
- If the KES released k_{d0}^A , Bob can use the `KeyUpdate` function to sequentially derive all subsequent keys $(k_{d1}^A, \dots, k_{di}^A)$ up to the state he claimed. He can then recover Alice's signature for state i and close the channel. This mechanism also acts as a punishment: if Alice was the malicious party trying to force an old state, Bob can now use her released initial key to derive the decryption key for any **later** state $j > i$ that is favorable to him, and close the channel at that state, stealing her funds.

Grease deviates from this original design in several ways:

- No transactions are shared with the KES. These are replaced with simple digitally-signed state messages.
- The protocol is generalized with the introduction of the channel id. This permits all parties to handle multiple concurrent channels.
- The dispute process uses a simple update count to significantly reduce the computational load on the KES.

4.8. Force-closing in Grease

4.8.1. Summary

Either party in the channel, the *claimant*, can trigger a force-close at any time. In short, the claimant informs the KES about the intent to force close, along with the current update count of the channel, n .

If the claim is valid, the KES starts the dispute window. If the counterparty, the *defendant*, has not presented a dispute proof before the window closes, the KES either

- encrypts and sends the defendant's ω_0 to the claimant (Option I).
- calculates the signature offset, wn_n and sends it to the claimant (Option II).

The data the claimant receives is a configuration option and is specific to the KES. In general, if the VCOF is very fast and cheap for the KES, then consider Option II. Otherwise, default to Option I. The biggest benefit of Option I combined with the dispute resolution improvement of the Grease protocol over AuxChannel is that the KES can be completely ignorant of the VCOF.

Important

The Grease design allows the KES to be completely independent and ignorant of the VCOF.

This is extremely useful if generating ZK-SNARKS on the KES platform is impossible, or very expensive. It also allows us to upgrade the VCOF on channels without having to change KES executables, which on blockchains is a non-trivial exercise.

A more detailed description of the process is given in the following sections.

```

onForceCloseRequestEvent(req: ForceCloseRequest) :
1  Validates the signature in req .
2  if valid:
3      Fetch the OpenChannel (Section 4.6.3) record with id = req.id => channel
        (Only allow channel parties to initiate a force-close)
4      if exists(channel) && req.Pc in [channel.Pa, channel.Pb] :
5          Create a new PendingChannelClose entry with status = Pending .
6          try to send a copy of the PendingChannelClose to the counterparty.
7          return Ok( ForceCloseResponse )
8      else
9          | return Fail( FailReason )
10 else
11 | return Fail( InvalidSignature )

```

Figure 8: The KES `HandleForceCloseRequest` algorithm.

4.8.2. Force-close procedure

If a claimant wishes to initiate a force close they send a `ForceCloseRequest` message to the KES which triggers an `onForceCloseRequestEvent` event. The KES responds to the request event with the `HandleForceCloseRequest` procedure, outlined in Figure 8.

The claimant can follow up with a `ClaimChannelRequest` message after the dispute window has closed. When the KES receives a `ClaimChannelRequest`, it triggers an `OnClaimChannelEvent` event.

The KES responds to this event with the `HandleClaimChannelRequest` algorithm, Figure 9.

Finally when the KES receives an `AckChannelClaimedEvent(channel_id)` signed by P_c it can permanently delete the corresponding `PendingChannelClose` record. It can also perform any other cleanup relating to the channel. At this point, the channel is closed and no trace of it remains on the KES.

4.8.3. ForceCloseRequest message

A party sends a `ForceCloseRequest` message to initiate the force-close process.

```

pub struct ForceCloseRequest<Curve> {
    /// The globally unique channel id
    channel_id: ChannelId,
    /// The public key of the party initiating the force close, $P_c$.
    claimant: PublicKey,
    /// The public key of the counterparty, $P_d$,
    defendant: PublicKey
    /// The most recent claimed state update count
    update_count_claimed: u64,
    /// A signature under $P_a$ for `id | update_count_claimed | Pc | Pd`
    signature: SchnorrSignature,
}

```

```

OnClaimChannelEvent(req: ClaimChannelRequest) :

1 Validates the signature in req.
2 if valid:
3   Fetch PendingChannelClose where id = req.id => pending
4   if !exists(pending) return Fail(No such channel)
5   if req.claimant ≠ pending.req.claimant return Fail(Unauthorized)
6   if current time ≥ pending.expiry
7     pending.req.defendant ⇒  $P_d$ 
8     pending.req.claimant ⇒  $P_c$ 
9     using Figure 5, decrypt  $\chi$  for pending.offsets where public key is  $P_c$  and ⇒  $\omega_0$ .
10    Case claim option:
11      Option I:  $\omega_0 \Rightarrow wn$ 
12      Option II: using Section 3.6.1 for n = pending.req.update_count_claimed, calculate
13       $wn_n \Rightarrow wn$ 
14      encrypt  $wn$  for  $P_c \Rightarrow \chi_c$  using Figure 4
15      update pending, set status to ForceClosed, add ( $P_c$ ,  $\chi_c$ , n) to chi
16      try send  $\chi_c$  to claimant
17 else return Fail("Dispute window still open")
18 else return Fail(InvalidSignature)

```

Figure 9: The KES `HandleClaimChannelRequest` algorithm.

4.8.4. PendingChannelClose message

The KES must track all channel that are under dispute.

```

pub struct PendingChannelClose {
  /// The globally unique channel id
  channel_id: ChannelId,
  /// A timestamp for when the dispute window closes
  expiry: u64,
  /// Initial offsets (encrypted)
  chi0: Vec<(PublicKey, EncryptedOffset)>
  /// A copy of the close request that intitiated the dispute
  req: ForceCloseRequest
  /// The status of the force-close channel
  status: PendingCloseStatus
  /// Encrypted offsets available for parties to close the channel
  chi: Vec<(PublicKey, EncryptedOffset, UpdateCount)>
}

pub enum PendingCloseStatus {
  /// The channel is pending force-closure and the dispute window is open.
  Pending,
  /// The dispute window is closed, and the channel is claimable by the claimant
  Claimable,
  /// The channel is abandoned. It is now claimable by either party
  Abandoned,
  /// The channel was closed via consensus and X value(s) are available for download.
  ConsensusClosed,
  /// The channel was claimed via force close,
  ForceClosed,
  /// The channel was claimed after being abandoned
}

```

```

    AbandonedClaimed,
    /// The channel was closed after a successful dispute
    DisputeSuccessful,
}

```

4.8.5. ClaimChannelRequest message

Once the dispute window has closed a claimant can claim their counterpart's ω_0 with a `ClaimChannelRequest` message:

```

pub struct ClaimChannelRequest {
    /// The globally unique channel id
    channel_id: ChannelId,
    /// The public key of the claimant, `P_c`
    claimant: PublicKey,
    /// A signature under $P_c$ for `id | Pc`
    signature: SchnorrSignature
}

```

The KES responds to a `ClaimChannelRequest` message by triggering an `onClaimChannelRequestEvent` event. It is handled by `handleClaimChannelRequest` as described in Figure 9.

A `ClaimAbandonedChannelRequest` message is identical to `ClaimChannelRequest`. However, its semantics are different:

- Either claimant **or** defendant may sign a `ClaimAbandonedChannelRequest` message.
- The KES will only allow the claim if `current_time ≥ pending.expiry + req.dw`, i.e. typically an additional 24 hours after the dispute window closes.
- After a successful claim, `pending.status` is set to `AbandonedClaimed`.


This message handles the case where a merchant has initiated a force-close but then has stopped responding. Without the ability to claim an abandoned channel, Alice has no way to recover her funds.

In summary, the various windows during the dispute phase are:

TIME	VALID SIGNERS	VALID MESSAGES
$t_{\text{close}} \rightarrow$	No-one. Dispute window	
$t_{\text{close}} + dw \rightarrow$	Claimant	<code>ClaimChannelRequest</code>
$t_{\text{close}} + 2dw \rightarrow t_D$	Either	<code>ClaimChannelRequest</code> (Claimant) <code>ClaimAbandonedChannelRequest</code> (Either)
$t_D \rightarrow \infty$	No-one. Record deleted	

Note

In the “abandoned” window, that is, between the closing of the dispute window and the record deletion (Section 4.10), the claimant can claim the channel using either type of message, but it is recommended to use `ClaimChannelRequest`.



Listing 7: The channel dispute sequence

4.9. Channel Dispute

Once a party (the defendant) becomes aware that there is a force-close in progress on one of her channels, she can:

1. Dispute the force close, and submit the data that would have closed the channel co-operatively,
2. Dispute the force close, proving a more recent channel state than the force-close state,
3. Do nothing and let the dispute window close.

4.9.1. Force close with consensus state

If Alice cannot prove a more recent state, under Option II, she can choose to do nothing at no cost to herself.

Under Option I, if there exists a prior channel state that is more disadvantageous to her than the claimed closing state, Alice is incentivized to publish a `ConsensusCloseRequest` (Section 4.9.1.1) which will close the channel in the state submitted by Bob in the `ForceCloseRequest`.

Note

It is strictly more expensive and time-consuming for both parties to close a channel under consensus rather than a co-operative close (Section 3.7), with exactly the same result.

When the KES receives a `ConsensusCloseRequest` message, it triggers an `onConsensusCloseRequestEvent` which leads to the `handleConsensusCloseRequest` methods carrying out Figure 10.

4.9.1.1. `ConsensusCloseRequest` message

The defendant submits a signed χ_n to allow the claimant to complete the refund transaction under the `update_count_claimed` in their `ForceCloseRequest`.

```
pub struct ConsensusCloseRequest {
    /// The globally unique channel id
    channel_id: ChannelId,
    /// The public key of the party initiating the force close, `P_c`
    claimant: PublicKey,
    /// The public key of the counterparty, $P_d$,
    defendant: PublicKey
    /// The most recent claimed state update count
    update_count_claimed: u64,
    /// The encrypted offset for state n,  $\chi_n$ 
    encrypted_offset: Scalar,
    /// A signature under $P_d$ for `id | update_count_claimed | Pa | Pb`
    signature: SchnorrSignature,
}
```

```
handleConsensusCloseRequest(req: ConsensusCloseRequest)
```

```
1 if !req.signature valid, return Fail( Unauthorized )
2 if exists(req.channel_id) in PendingChannelClose records as p:
3   if p.defendant  $\neq$  req.defendant OR p.claimant  $\neq$  req.claimant, return Fail( Invalid )
4   if p.update_count_claimed  $\neq$  req.update_count_claimed, return Fail( Invalid )
5   Decrypt req.encrypted_offset  $\rightarrow$   $w_n$  (Figure 5), else return Fail( Invalid )
6   Encrypt  $w_n \rightarrow \chi_n^D$  under  $P_D$  (Figure 4)
7   Update p: set p.chi_D =  $\chi_n^D$ , p.status = ConsensusClosed
8   Try send  $\chi_n^D$  to defendant.
9 else return Fail( NotFound )
```

Figure 10: Resolving a channel force-close via consensus

4.9.2. Dispute with recent state proof

The `ForceCloseRequest_message`[†] contains a signature from both parties signing the update count. This is the same information that parties exchange on *every* channel update. Therefore, it is now very straightforward for the KES to verify a dispute.

1. Alice provides a `DisputeChannelState` message with an `UpdateRecord` that has been signed by Bob:

```
pub struct DisputeChannelState {
    /// The globally unique channel id
    channel_id: ChannelId,
    /// The public key of the claimant, $P_c$
    claimant: PublicKey,
    /// The public key of the defendant, $P_d$,
    defendant: PublicKey
    /// The channel update count
    update_count: u64,
    /// Other data not relevant to the KES, but included in the signed message
    ...
    /// A signature under $P_b$ for `id | update_count | Pa | Pb | ...`
    update_record: SchnorrSignature,
    /// This record signed by the defendant
    signature: SchnorrSignature,
}
```

The KES makes the following checks:

1. `update_count > n_claimed`, i.e. that Alice's proof is for a more recent state than Bob's force-close request.
2. P_a and P_b are identical in both the force-close request and dispute channel state messages.
3. The `signature` in the dispute channel state message is valid and signed by the defendant.
4. The `update_record` signature is valid and signed by the claimant.

If all checks pass, the KES decrypts ω_0^B with Figure 5, encrypts it to P_A using Figure 4 and sends the encrypted ω_0^B to Alice.

4.9.3. Do Nothing

Alice can do nothing, in which case the dispute window will eventually close and Bob can claim the missing signature offset via Figure 9.

4.10. Channel cleanup

4.10.1. After co-operative close

- The merchant will deliver proof that the channel has closed. This proof consists of a `ChannelClose` message corresponding to the channel with id `id` signed by each of the Merchant and Customer.
- If, and only if, both signatures are valid, the KES **must**
 - destroy all data associated with channel `id` insofar as it is possible.
 - return any resources (e.g. deposits or collateral) belonging to either party that were being held as part of the KES, net of any fees associated with maintaining the KES.

4.10.2. Hygiene

To maximize privacy, the KES will delete old channels after a reasonable storage period. t_D is configurable by the KES, but should typically be one month after the claim period starts, i.e.

$$t_D = \text{expiry} + 2 \text{ dw} + \text{one_month} \quad (4)$$

Periodically, the KES will scan all records in `PendingChannelClose` and

- delete all records where `current_time > t_D`
- where `current_time > expiry + 2dw`, set `status = Abandoned`
- where `current_time > expiry + dw`, set `status = Claimable`

4.11. Utility functions

The KES must also expose the following utility functions:

- `getChannelDefinition(channel_id, auth)`
- `getPendingClose(channel_id, auth)`

This allows the customer to query the channel lifecycle independently of the merchant. In both cases authorization consists of a Schnorr signature signing the channel id and requester's public key.

- If the signature is valid, but the channel record does not exist, the KES returns "Not found".
- If the signature is invalid, the KES returns "Unauthorized".
- If the channel record exists, **and** the signer's public key is recorded in the record, return the record. Otherwise return "Not found".

5. Trust assumptions in the KES

Section 4 mapped out the cryptography that the KES needs to manage. It did not cover how it might be implemented in a trust-minimizing manner. This section examines some practical approaches for implementing the KES as well as discussing some of the trade-offs involved.

A trustless (or trust-minimized) KES deployment must satisfy the following properties:

- **Collusion resistance.** The primary threat is the KES colluding with one party to share the counterparty's ω_0 outside of the dispute protocol. This would allow the colluding party to unilaterally drain the channel.
- **Secrecy.** Neither channel party can access their counterparty's ω_0 outside of the conditions defined by the protocol.
- **Conditional release.** ω_0 is released if and only if the conditions in the dispute and claim phases are satisfied.
- **Verifiability.** All participants must be able to verify that the KES is operating correctly, and only according to the prescribed protocols.
- **Censorship resistance.** in many respects, censorship is equivalent, or arguably, worse than collusion, since it can make it appear that a party has become unresponsive, and the other two parties can carry out a completely valid-looking force-close.

We can capture these requirements in these three properties:

1. We must be certain, and able to independently verify, that the code describing the KES algorithm, **and only that code** is executed as part of the KES' duties.
2. We must be able to verify, within a suitable confidence threshold, that no coalition of less than the security threshold can recover ω_0 (or anything from which ω_0 can be derived) outside of the conditions in the dispute and claim phases.
3. The KES must not be able to ignore valid messages from parties (censorship resistance).

Property 1 can be met to within a reasonable approximation by executing code on a decentralized public smart contract blockchain. The code does not have to be private.

Property 2 is more challenging. For trustlessness, we need to convince ourselves that the KES as a whole is *unable* to leak k_K to any other party within the constraints of a set of security assumptions. This is possible in a Trusted Execution Environment (), but at the time of writing, no smart contract system has a way of generating and then storing a verifiably random number in private storage (see).

Property 3 requires two ingredients that no centralized deployment can provide: (a) request submission and ordering must be observable by parties other than the KES operator, so that censorship is *detectable*, and (b) the entity that finalises which requests are processed must be a quorum rather than a single operator, so that censoring a request requires corrupting an honest majority of that quorum rather than a single machine. BFT consensus on a public blockchain delivers both: the mempool (or its equivalent) gives (a), and the validator set gives (b). Public blockchains are therefore the natural vehicle for property 3; centralized solutions are excluded.

5.1. The problem with private state on blockchains

There are several blockchain projects that purport private smart-contract support. Lead among them are Aleo and Aztec, both using ZK-SNARKs for private computation proofs. DarkFi is also in development, but is less mature. Tari supports confidential payments, and its architecture allows for private computation in future, but the current testnet design exposes all state to validators.

For brevity and expedience, the remainder of this discussion will assume Aztec as the chosen platform for the KES. Grease already makes use of Noir (Aztec’s smart-contract language) and there are no features in the other projects that Aztec doesn’t already have. Aztec is currently in “Alpha-mainnet” at the time of writing – the rollup is live on Ethereum Mainnet with 500 validators, but the team considers the platform an alpha release with several [critical bugs](#)[†].

In Aztec’s model, private state is encrypted and owned by users who hold the corresponding decryption keys. Modifying private state requires nullifying existing records and creating new ones (a UTXO model), which preserves privacy of state transitions. However, *someone always holds the decryption key* for any given piece of private state. In fact, there are no ZK blockchains at the time of writing that offer purely private storage.

Until this is the case, we are forced down the path of multiparty computation. k_K needs to be sharded so that property 2 is satisfied. No-one knows the whole secret but say, t -of- n parties can co-ordinate to sign messages. The Olaf protocol based on Frost3⁹ is one way of carrying this out, though there are others^{10–12}.

5.1.1. A hybrid on-chain off-chain design

Figure 11 illustrates the architecture for a hybrid on-chain off-chain design. KES state storage and signature aggregation and verification is maintained on-chain in smart contracts. The KES private key k_K is sharded among n signers that run “hot-wallet” style daemons to respond to signature requests. A threshold t is chosen such that any t signers can co-operate to sign a message with k_K . No single signer knows that key, and t must be large enough such that the probability of collusion is suitably low.

FrostCoordContract:

- Public functions on Aztec L2.
- Stores the signer registry and aggregate public key P_K .
- Verifies FROST signatures.
- Does not orchestrate DKG or signing rounds – those happen off-chain between committee members.
- The contract is a registry and verification endpoint.

KesStateContract:

- Public functions managing channel state.
- Stores `OpenChannel` and `PendingChannelClose` records in public storage.
- Enforces the dispute state machine using L1-derived timestamps.
- Issues release authorizations.

Signer daemons: (n instances)

- Off-chain processes.
- Each holds a shard of k_K in local memory.
- Watches the L2 public state for events.
- Participates in peer-to-peer DKG, threshold decryption, and FROST signing rounds. Submits results to the L2 contract via standard transactions.




Figure 11: KES hybrid-architecture sketch for Aztec

This design satisfies all three properties above, at the cost of much higher complexity. Liveness is likely to be on the order of 1-2 minutes for channel initialization, which while fairly slow, is still acceptable, since the channel won't truly be open until the monero funding transactions have been included in a block.

The choice of t is a direct trade-off: smaller t improves liveness but increases the chance of collusion. A reasonable starting point is $t \approx 2\frac{n}{3}$, matching the BFT threshold of the underlying L2.

FROST gives the signing committee a way to produce Schnorr signatures under P_K without ever reconstructing k_K . It does not cover recovering ω_0 from the encrypted offset χ that the parties submit. `DecryptMessage` (Figure 5) requires the holder of k_K to compute $k_K \cdot R$, which is a Diffie-Hellman operation, not a Schnorr signature. The hybrid design therefore provides a second threshold primitive in addition to FROST: threshold decryption of the hashed-ElGamal ciphertexts produced by `EncryptMessage` (Figure 4).

Fortunately, the encryption scheme in Figure 4 is friendly to this. Threshold decryption is possible using the same shares that parties obtained during the Olaf DKG ceremony.

Section 4.4 defines a per-channel keypair $k_g = \kappa \cdot k_K$ with $P_g = k_g \cdot G$, and the proof-of-knowledge signatures returned during channel opening are signed under k_g , not k_K . The hybrid design inherits this without needing a fresh DKG per channel because Shamir secret sharing is linear over the scalar field, the share

$$k_{g_i} = \kappa \cdot k_{K_i} \quad (5)$$

is a valid Shamir share of k_g with the same threshold t and the same evaluation points as the original sharing of k_K .

Section 4.4 requires the KES to discard κ immediately after deriving k_g . In the hybrid design this requirement is relaxed. Under the standing assumption that fewer than t signers collude, no-one learns κ so there is nothing to discard.

The added complexity, and by extension, cost, means that KES signers will probably need to be compensated for running daemons. As mentioned in Section 5.3, adding a fee component to the KES is relatively straightforward, but is out of scope for this white paper.

5.1.2. Additional applications and distribution of risk

One may observe that the design in Figure 11 could easily be generalized to store *any* encrypted data in a decentralised, censorship-resistant manner. To wit, the design migrates from a *Key* escrow service to a *Document* escrow service.

If one were to build this out, the `KesStateContract` is but one of many types of data the “DES” signers are prepared to encrypt. We will not dive down this rabbit hole here, but this is worth exploring further if one concludes that this design needs a funding mechanism to be sustainable, but that Grease alone may or may not provide sufficient volume to support the KES on its own.

Possible applications include:

- On-chain insurance contracts,
- Decryption of documents under certain conditions, e.g. crypto inheritance strategies with dead-man switches,
- On-chain wagers and prediction markets

5.2. Using TEEs

TEEs satisfy Properties 1 and 2 to a reasonable degree (see the caveats below). However, one cannot run a KES on a single TEE since a malicious KES operator could simply censor messages from say, Alice, which would allow Bob to submit a force-close request and extract Alice’s ω_0 after the time-out without the KES operator even needing to break the enclave boundary!

Therefore, in order to satisfy property 3, you need distributed TEE infrastructure with the ability to form BFT consensus across the network. As luck would have it, such a system already exists – the Secret Network¹³. By leveraging the infrastructure and consensus machinery already deployed in Secret, deploying a KES in a distributed TEE network becomes fairly straightforward.

5.2.1. The KES on the Secret network

Deployment to the Secret Network looks almost identical to a centralized KES contract (Figure 12) – the network is purely used to provide censorship resistance.




Figure 12: KES contract outline for the Secret Network

Every node must create the same private key in order to come to consensus. The Secret network actually provides a secret random number to enclaves for each block, so the key can be derived without any external observer knowing what it is, for example

```
// Derive k_K from consensus-seeded randomness.
let seed = ctx.env.block.random.ok_or(Error::NoRandomness)?;
let mut okm = [0u8; 64];
Hkdf::<Sha512>::new(Some(b"kes.contract.master_key.v1"), seed.as_slice())
    .expand(b"k_K", &mut okm)
    .unwrap();
let k_k = Scalar::from_bytes_mod_order_wide(&okm);
let p_k = &k_k * BASEPOINT;
```

5.2.1.1. Advantages

- Simple deployment – almost identical to a centralized KES.
- Performance is very good. Tendermint BFT means finality in one block. Block intervals are 5-6s.
- Permissionless operation
- Liveness and availability guarantees are identical to that of the the Secret network.
- Secret contracts are pure Rust compiled into WASM. Thus Curve25519 can be used for $\mathbb{E}_K$, making the DLEQ proofs trivial.
- No human ever sees the private key.

5.2.1.2. Disadvantages

- **TEE single point of trust.** Until Secret diversifies its hardware requirements, if SGX is compromised network-wide, all key shares for all signer contracts are exposed simultaneously.
- **Limited decentralisation.** 80 validators, all running the same Intel SGX hardware. Although, for censorship resistance, this is likely sufficient.

5.2.1.3. Intel SGX exploits

The locus of trust for a Secret-based KES shifts to the TEE hardware.

Currently Secret requires the use of Intel's SGX trusted enclaves, although they are starting to support other hardware¹⁴ including AMD SEV-SNP.

Unfortunately, SGX has a fairly long history of vulnerabilities.

Pre-2024, when Secret was fully permissionless (for validators), the Foreshadow (2018), SGAXe (2020) and Downfall exploits may have led to the private key leaking had the KES been running. Post-2024, Secret validators go through a vetting process of sorts. Even though this gating is friction for decentralization, it has improved the security of the Secret network. Several exploits that involved faking attestation certificates, including *wiretap.fail*, would have been caught by the Secret governance process.

Overall, for Grease, the conclusion is that while Secret's security is improving, SGX is still its Achilles' heel. That said, given how simple the deployment is, it's the author's opinion that a Secret-based KES would be fine for low- to medium value channels as a stop-gap while a more robust KES is built out, rather than having no KES at all, or only trusted centralized KESs (recall that Grease supports multiple KESs— selecting the KES is part of the channel opening negotiation).

5.3. A note on economic incentives for centralized KEServices.

We also considered centralized KES systems that are *incentivized* to behave correctly.

The idea is that the KES charges a small fee for escrowing the ω_0 values. In turn, the KES stakes their reputation, or even monetary tokens on being honest.

Fees are a straightforward extension to Grease:

- There are already multiple rounds of communication between Customer, Merchant and KES during channel establishment. A KES fee could be

grafted into the existing messages.

- The merchant is already responsible for communicating with the KES. He would be responsible for paying the fee, and recovering it by

including it into the channel payments. If the KES' currency is not Monero, the merchant can manage this as well by aggregating conversion over multiple channel openings.

Let's assume that fraud proofs are relatively easy to produce in Grease, since both parties to the channel keep signed records of every channel update. If a channel was closed with an outdated state, the aggrieved party should be able to demonstrate this convincingly.

However, there is a problem: Alice and Bob can collude *against the KES*:

- Alice gives ω_0^A to Bob directly (off-protocol, without KES involvement).
- Bob closes the channel on Monero.
- Alice files a fraud proof, claiming the KES colluded.
- The KES stake is slashed / or the KES' reputation is destroyed.

There is no straightforward design that mitigates this attack. For example, requiring Alice and Bob to escrow additional funds that could be slashed by a counterproof breaks the very point of Grease in the first place: Simple, low-friction payment channels.

If preventing collusion against the KES is not practical, then no fraud-proof-based system is viable, which essentially excludes centralized KES implementations from all but very low-stakes applications (e.g. conferences, or festivals, where the channels are naturally short-lived and carry only a few tens of dollars at most).

5.4. Summary

Centralized KES deployments are simplest and fastest, but they require parties to trust that the KES won't collude with their counterparty to steal funds. This is fine for testing and low-stakes applications, but is not sufficient for Grease's broader adoption goals.

The Secret Network provides a BFT network of Trusted Execution Environments (TEE) that –on the face of it– deliver exactly what the KES needs: private secret storage, verifiable computation and censorship resistance. Unfortunately the “T” in TEE is not particularly well-earned. Nevertheless, the Secret network is a decent interim solution for the KES, providing a bridge between a purely centralized KES and a longer-term SNARK-based solution.

A blockchain-only trustless KES deployment is not possible with current ZK-SNARK based platforms, since they cannot store and use private state. Aztec, as arguably the leading ZK blockchain, has private execution on *client devices*. The KES secret can't be placed in the hands of any single entity, since that invites collusion. Therefore a hybrid model is posited: KES coordination and storage on-chain coupled with an Olaf/FROST3 distributed signing committee. Individual signers operate as an always-on network service.

The proposed approach forward for Grease would be the following:

1. Start with a centralized KES for the beta phase. This allows the UX of the channel parties to be ironed out and the core cryptography of the escrow service to be battle-tested.
2. Deploy the first production KES on the Secret network. We don't particularly like Intel SGX, but it's a relatively cheap deployment, is better than a purely centralized KES, and bridges the gap until,
3. A hybrid escrow service on Aztec is complete; or if private storage is solved, simply migrate off of TEEs to a SNARK-based platform. In both cases, there may be a market for providing a more generalized “Document Escrow Service”, though determining the level of demand for such a service is outside the scope of this document.

6. Future extensions and known limitations

The Grease protocol represents the first attempt to extend the high-security features of Monero while also using the problem-solving flexibility of the latest Turing-complete ZKP tools. Given the rate at which the ZKP technology is advancing there may be many more opportunities to extend Monero's security to new features and markets, connecting the future of Monero with the larger blockchain community.

Table of Algorithms

7. Nomenclature

7.1. Symbols

SYMBOL	DESCRIPTION
$\mathbb{G}_2$	Generator point for curve Baby JubJub
G_1	Generator point for curve Ed25519
L_2	The prime order of curve Baby JubJub
L_1	The prime order for curve Ed25519
ω_i	The witness value for party i
T_i	The public point corresponding to ω_i on curve Baby JubJub
S_i	The public point corresponding to ω_i on curve Ed25519

7.2. Subscripts

SUBSCRIPT	REFERENT
BJJ	Curve Baby JubJub
Ed	Curve Ed25519
M	The peer playing the role of merchant
C	The peer playing the role of customer
Initiator	The peer playing the role of initiator
Responder	The peer playing the role of responder

References

1. What is Monero?. <https://www.getmonero.org/get-started/what-is-monero/>^o
2. Chainalysis. All About Monero. <https://www.chainalysis.com/blog/all-about-monero/>^o
3. Alexander Culafi. Monero and the complicated world of privacy coins. January 24, 2022. <https://www.techtarget.com/searchsecurity/news/252512394/Monero-and-the-complicated-world-of-privacy-coins>^o
4. Zhimei Sui, Joseph K. Liu, Jiangshan Yu, Man Ho Au, Jia Liu. *Auxchannel: Enabling Efficient Bi-Directional Channel for Scriptless Blockchains.*; 2022. <https://eprint.iacr.org/2022/117.pdf>^o
5. Sui Z, Liu JK, Yu J, Qin X. *Monet: A Fast Payment Channel Network for Scriptless Cryptocurrency Monero.*; 2022. <https://eprint.iacr.org/2022/744.pdf>^o
6. CJ, Sean Coughlin. Grease: A minimal Monero Payment Channel implementation for Monero. <https://cfp.twed.org/mk5/talk/QYDGPM/>^o
7. jeffro256. Personal communication: Signature changes post-FCMP. Published online June 22, 2025.
8. Daniel J. Bernstein. Curve25519: new Diffie-Hellman speed records. February 9, 2006. <https://cr.yp.to/ecdh/curve25519-20060209.pdf>^o
9. Hien Chu, Paul Gerhart, Tim Ruffing, Dominique Schröder. *Practical Schnorr Threshold Signatures Without the Algebraic Group Model.*; 2023. <https://eprint.iacr.org/2023/899.pdf>^o
10. Chelsea Komlo, Ian Goldberg. *FROST: Flexible Round-Optimized Schnorr Threshold Signatures.*; 2020. <https://eprint.iacr.org/2020/852.pdf>^o
11. Tim Ruffing, Jonas Nick. ChillDKG: Distributed Key Generation for FROST (BIP draft). 2024. <https://github.com/BlockstreamResearch/bip-frost-dkg/blob/master/README.md>^o
12. Rosario Gennaro, Stanislaw Jarecki, Hugo Krawczyk, Tal Rabin. Secure Distributed Key Generation for Discrete-Log Based Cryptosystems. *Journal of Cryptology*. 2007;20(1):51-83. <https://link.springer.com/content/pdf/10.1007/s00145-006-0347-3.pdf>^o
13. The Secret Network. <https://scrt.network/>^o
14. Secret Labs. SecretVM: Expanding Confidential Computing with AMD SEV-SNP. Secret Network Blog. February 11, 2026. Accessed April 9, 2026. <https://scrt.network/blog/secretvm-expanding-confidential-computing-with-amd-sev-snp>^o

Index of Figures

Figure 1	Grease VCOF Update function	18
Figure 2	The ZK-SNARK Grease <code>VCOFProve</code> function	18
Figure 3	Grease VCOF Verify function	18
Figure 4	The <code>EncryptMessage</code> algorithm.	23
Figure 5	The <code>DecryptMessage</code> algorithm.	23
Figure 6	Diffie-Hellman shared secret derivation for new channels	23
Figure 7	The KES validates each new channel request, returning proofs of knowledge if successful.	26
Figure 8	The KES <code>HandleForceCloseRequest</code> algorithm.	29
Figure 9	The KES <code>HandleClaimChannelRequest</code> algorithm.	29
Figure 10	Resolving a channel force-close via consensus	34
Figure 11	KES hybrid-architecture sketch for Aztec	38
Figure 12	KES contract outline for the Secret Network	40

Index of Tables

Table 1	The role of the KES during the channel lifecycle	22
Table 2	Equivalent curve points in grease	27

Index of Listings

Listing 1	The New Channel proposal sequence	10
Listing 2	Establishing a new Channel	12
Listing 3	Establishing a new Channel, continued	13
Listing 4	Creating a new multisig wallet	15
Listing 5	Updating a Channel	17
Listing 6	The Channel Close sequence	21
Listing 7	The channel dispute sequence	32